CSE 332 Winter 2026 Lecture 16: Graphs 3

Nathan Brunelle
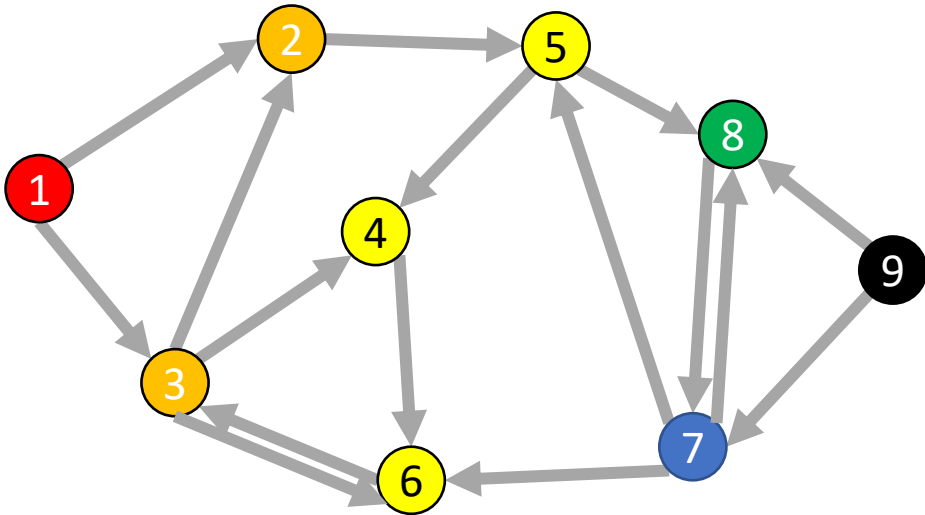
http://www.cs.uw.edu/332

Breadth-First Search

- Input: a node s

- Behavior: Start with node s, visit all neighbors of s, then all neighbors of neighbors of s, …

- Visits every node reachable from s in order of distance

- Output:
  - How long is the shortest path?
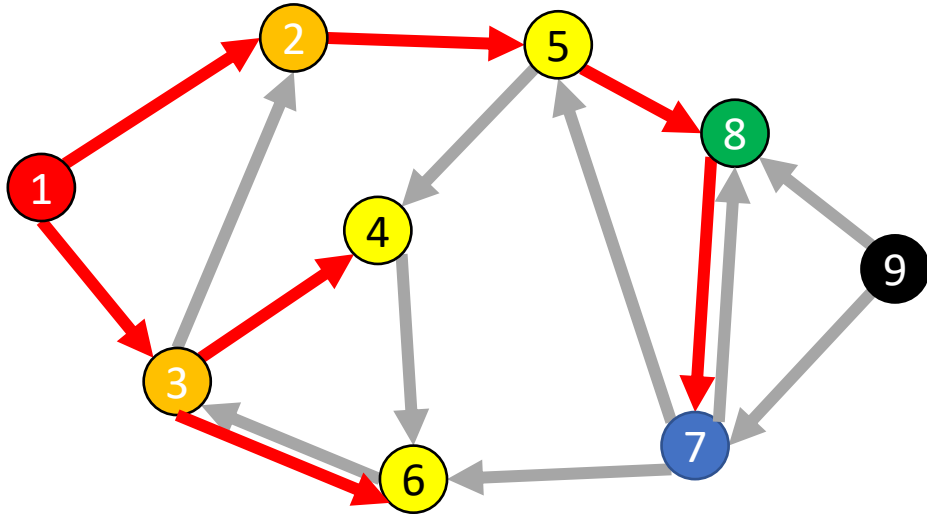 - Is the graph connected?

BFS

Running time: $\Theta(|V| + |E|)$

```
void bfs(graph, s){
    found = new Queue();
    found.enqueue(s);
    mark s as "visited";
    While (!found.isEmpty()){
        current = found.dequeue();
        for (v : neighbors(current)){
            if (! v marked "visited"){
                mark v as "visited";
                found.enqueue(v);
            }
        }
    }
}
```
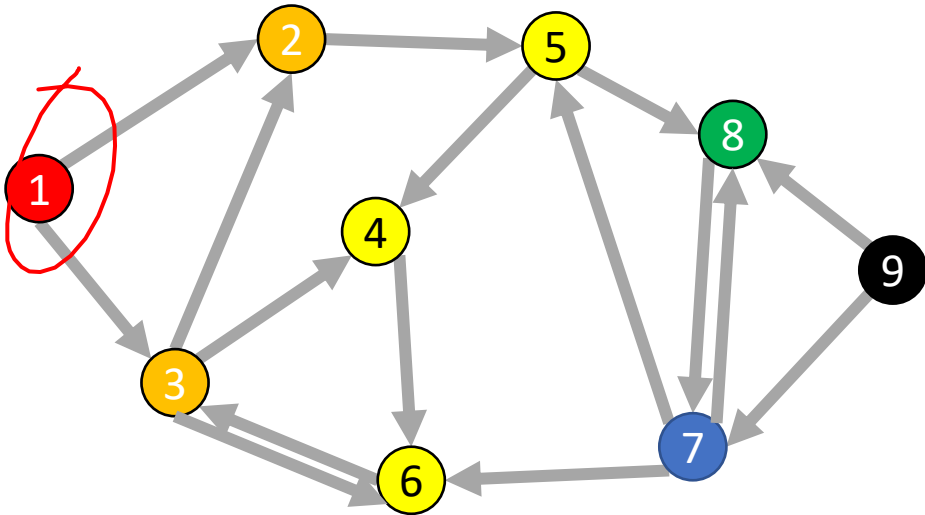
Find Distance (unweighted)

Idea: when it's seen, remember its "layer" depth!

```
int findDistance(graph, s, t){
        found = new Queue();
        layer = 0;
        depth of s = 0;
        found.enqueue(s);
        mark s as "visited";
        While (!found.isEmpty()){
                current = found.dequeue();
                layer = depth of current;
                for (v : neighbors(current)){
                        if (! v marked "visited"){
                                mark v as "visited";
                                depth of v = layer + 1;
                                found.enqueue(v);
                        }
                }
        }
        return depth of t;
}
```

Find Distance – Worked Example

Node	Visited?	Depth
1	T	0
2	T	1
3	T	1
4	T	2
5	T	2
6	T	2
7		
8	T	3
9		

For each node:
  update current layer
For each unvisited neighbor:
  add that neighbor to a queue
  mark that neighbor as visited
  set neighbor's layer to be current layer + 1

Queue: 1 2 3 5 4 6 8

Shortest Path - Idea

For each node:
 For each unvisited neighbor:
  add that neighbor to a queue
  mark that neighbor as visited
  set neighbor's previous to be the current node

1,2,5,8

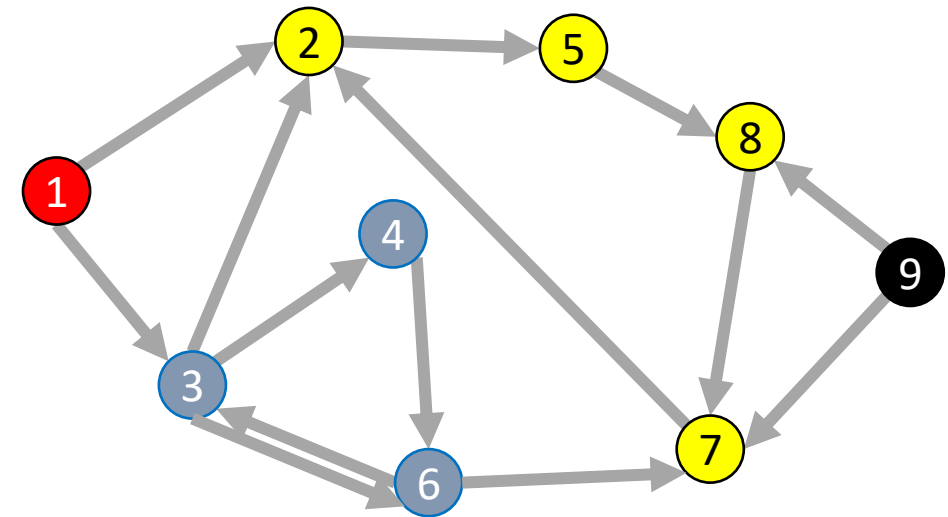| Node	Visited?	Previous
1	T	—
2	T	1
3	T	1
4	T	3
5	T	2
6	T	3
7		
8	T	5
9		

Queue: 2 3 5 4 6 8

Depth-First Search

- Input: a node s

- Behavior: Start with node s, visit one neighbor of s, then all nodes reachable from that neighbor of s, then another neighbor of s,…
  - Before moving on to the second neighbor of s, visit everything reachable from the first neighbor of s

- Output:
  - Does the graph have a cycle?
  - A **topological sort** of the graph.
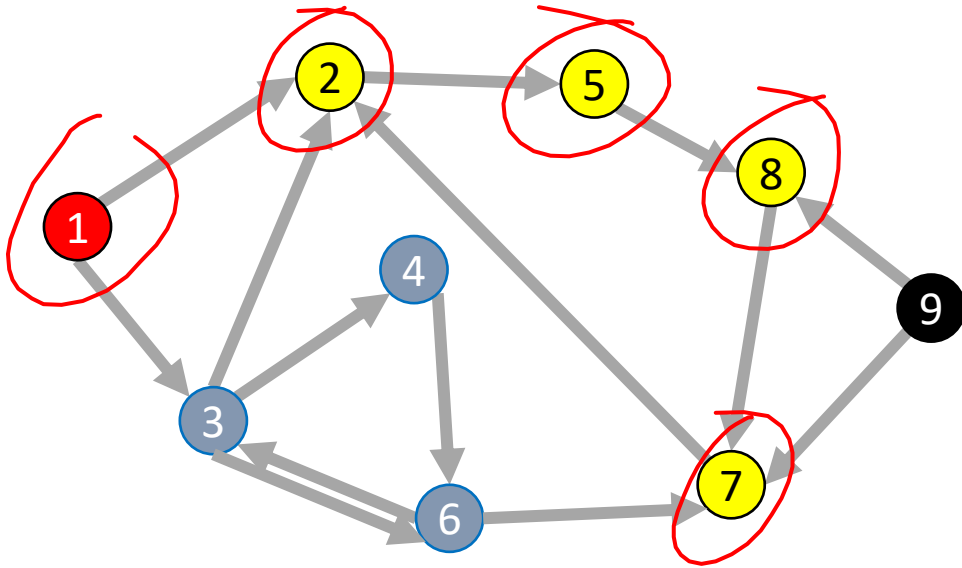
DFS Recursively (more common)

```
void dfs(graph, curr){
    mark curr as "visited";
    for (v : neighbors(current)){
        if (! v marked "visited"){
            dfs(graph, v);
        }
    }
    mark curr as "done";
}
```

DFS – Worked Example

Node	Visited?	Done?	Other Info
1	✓		
2	✓	✓	
3	✓		
4			
5	✓	✓	
6			
7	✓	✓	
8	✓	✓	
9			

Starting from the current node:
for each unvisited neighbor:
  mark the neighbor as visited
  do a DFS from the neighbor
mark the current node as done

(Call) Stack: 1 1 3

Using DFS

- Consider the "visited times" and "done times"
- Edges can be categorized:
  - Tree Edge
    - (a, b) was followed when pushing
    - (a, b) when b was unvisited when we were at a
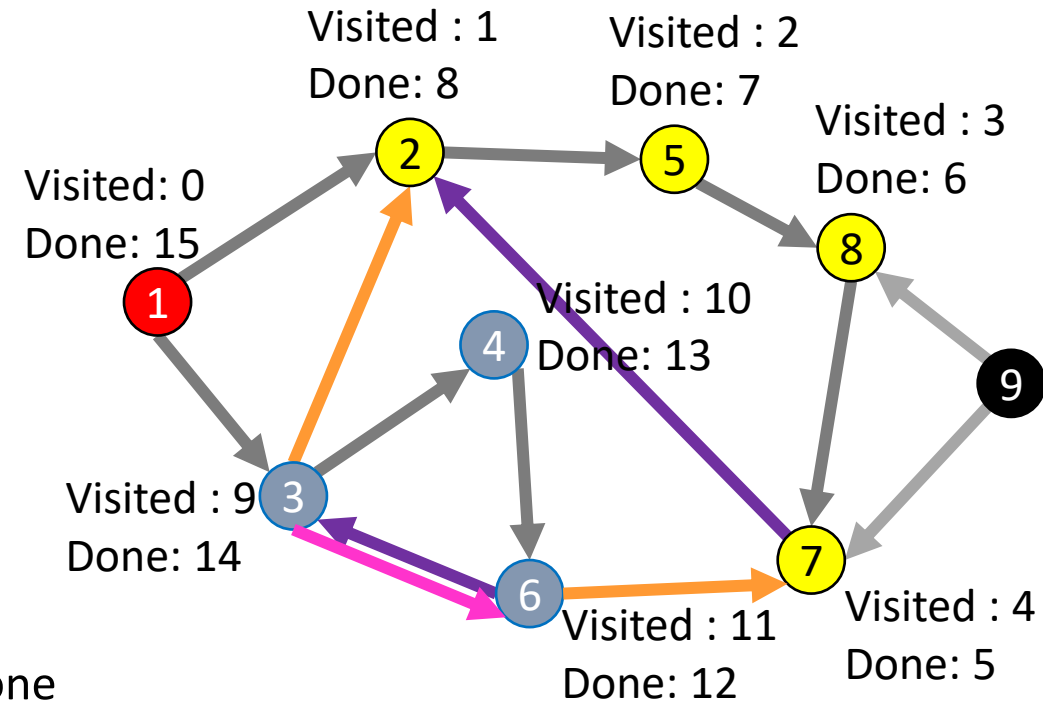 - Back Edge
    - (a, b) goes to an "ancestor"
    - a and b visited but not done when we saw (a, b)
    - $t_{visited}(b) < t_{visited}(a) < t_{done}(a) < t_{done}(b)$
  - Forward Edge
    - (a, b) goes to a "descendent"
    - b was visited and done between when a was visited and done
    - $t_{visited}(a) < t_{visited}(b) < t_{done}(b) < t_{done}(a)$
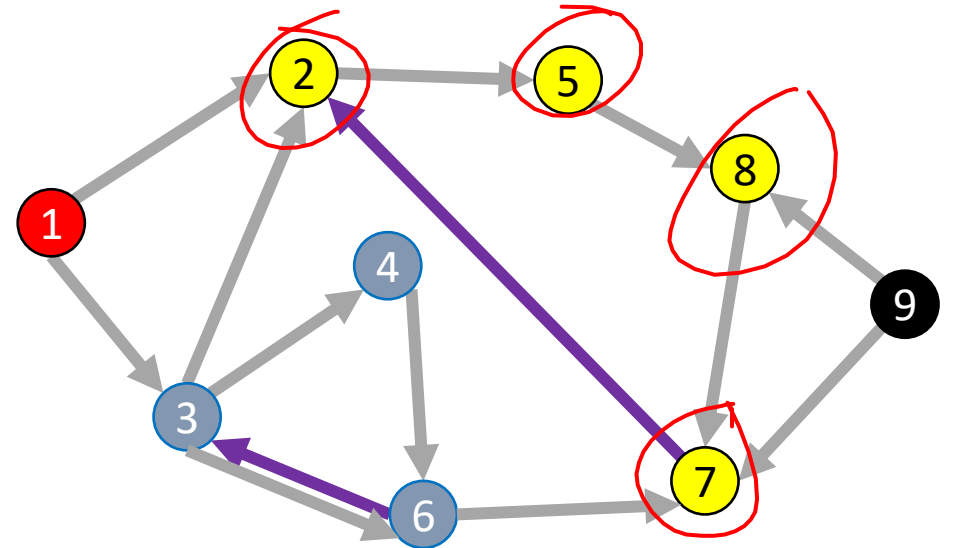 - Cross Edge
    - (a, b) goes to a node that doesn't connect to a
    - b was seen and done before a was ever visited
    - $t_{done}(b) < t_{visited}(a)$

Visited : 1
Done: 8

Visited : 2
Done: 7

Visited : 3
Done: 6

Visited: 0
Done: 15

Visited : 10
Done: 13

Visited : 9
Done: 14
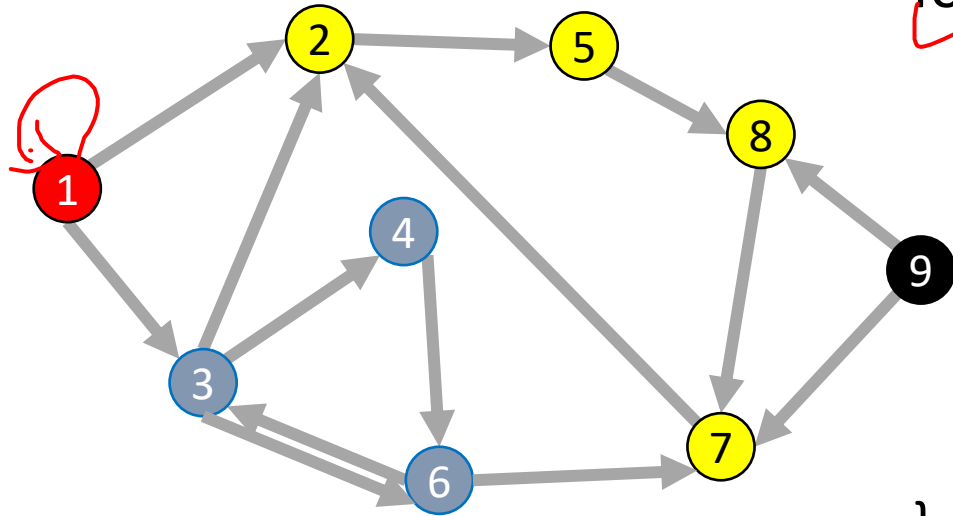
Visited : 11
Done: 12

Visited : 4
Done: 5

Back Edges

- Behavior of DFS:
  - "Visit everything reachable from the current node before going back"

- Back Edge:
  - The current node's neighbor is an "in progress" node
  - Since that other node is "in progress", the current node is reachable from it
  - The back edge is a path to that other node
  - **Cycle!**

Cycle Detection

```
boolean hasCycle(graph, curr){
        mark curr as "visited";
        cycleFound = false;
        for (v : neighbors(current)){
                if (v marked "visited" && ! v marked "done"){
                        cycleFound=true;
                }
                if (! v marked "visited" && !cycleFound){
                        cycleFound = hasCycle(graph, v);
                }
        }
        mark curr as "done";
        return cycleFound;
}
```

12

Cycle Detection – Worked Example

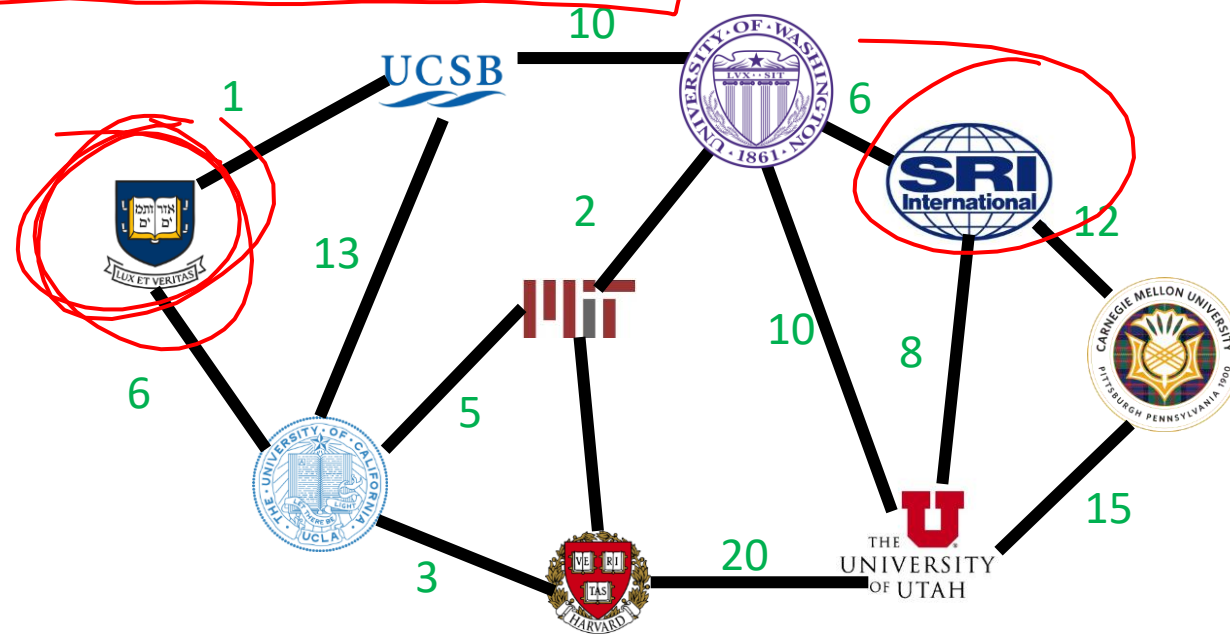| Node	Visited?	Done?	Other Info
1			
2			
3			
4			
5			
6			
7			
8			
9			

Starting from the current node:
for each non-done neighbor:
if the neighbor is visited:
we found a cycle!
else:
mark the neighbor as visited
do a DFS from the neighbor
mark the current node as done

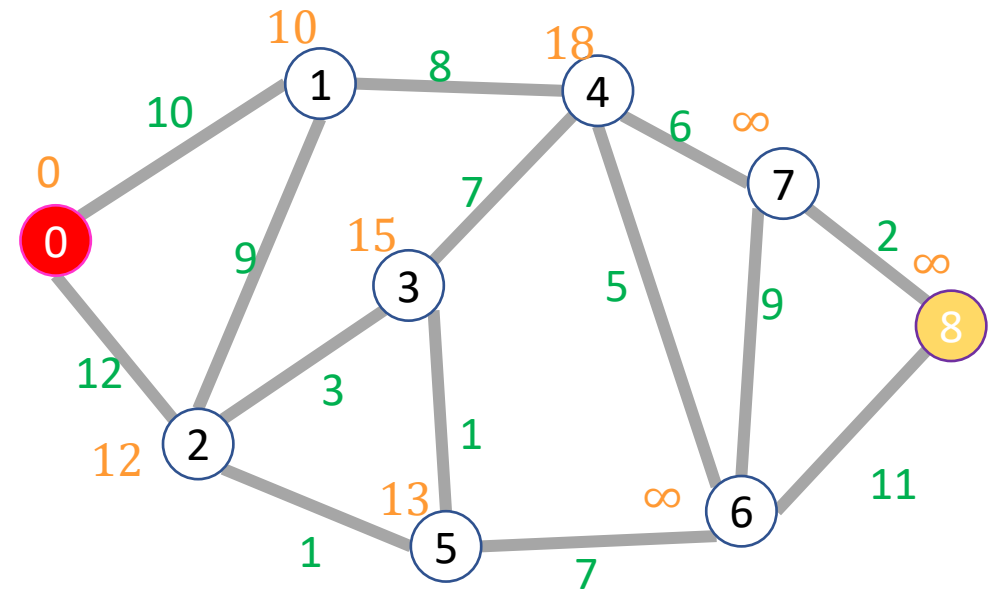(Call) Stack:

Single-Source Shortest Path

Find the quickest way to get from UVA to each of these other places

Given a graph $G = (V, E)$ and a start node $s \in V$, for each $v \in V$ find the least-weight path from $s \to v$ (call this weight $\delta(s, v)$)

(assumption: all edge weights are positive)

Dijkstra's Algorithm

- Input: graph with **no negative edge weights**, start node s, optional end node t

- Behavior: Start with node s, repeatedly go to the incomplete node "nearest" to s

- Output:
  - Distance from start to end
  - Distance from start to every node

Dijkstra's

Distance to start = 0
Add the start node to PQ with priority 0
Mark start as "seen"
While the PQ is not empty:
    curr = PQ.extract();
    mark curr as "done"
    for each neighbor v of curr:
        d = distance to curr + weight of (curr,v)
        if v is not "seen":
            mark v as "seen"
            distance to v = d
            PQ.add(v, d);
        if v is not "done" && d < distance to v:
            distance to v = d
            PQ.decreaseKey(v, d)

Loops $|E|$ times

Worst case $\Theta(\log|V|)$ each

Idea: When a node is the closest not-done thing to the start, we have found its shortest path

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

Seen = added to the priority queue
Done = removed from the priority queue
When it's done we've found the shortest path to that node

Running time: $\Theta(|E|\log|V|)$

16

Dijkstra's Algorithm

Start: 0

End: 8

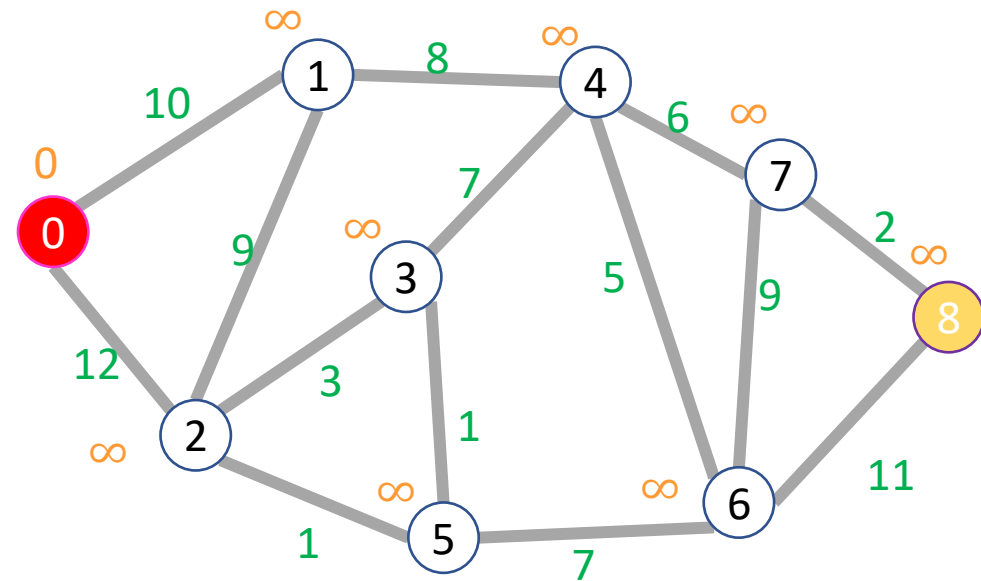| Node	Seen?	Done?	Distance
0	T	F	0
1	F	F	∞
2	F	F	∞
3	F	F	∞
4	F	F	∞
5	F	F	∞
6	F	F	∞
7	F	F	∞
8	F	F	∞

Idea: When a node is the closest not-done thing to the start, we have found its shortest path

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

Dijkstra's Algorithm

Start: 0

End: 8

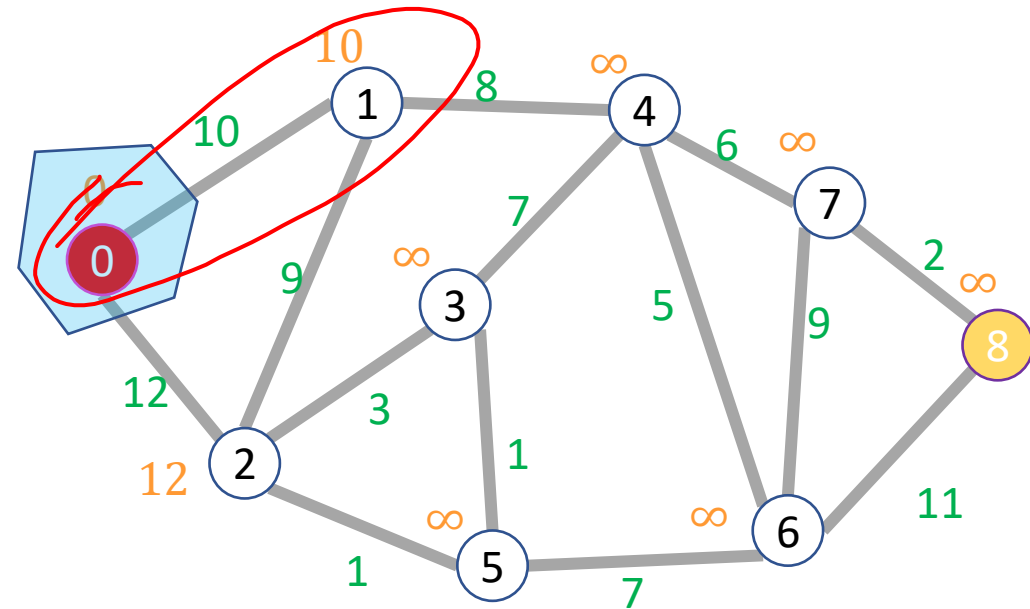| Node	Seen?	Done?	Distance
0	T	T	0
1	T	F	10
2	T	F	12
3	F	F	∞
4	F	F	∞
5	F	F	∞
6	F	F	∞
7	F	F	∞
8	F	F	∞

Idea: When a node is the closest not-done thing to the start, we have found its shortest path

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

18

Dijkstra's Algorithm

Start: 0

End: 8

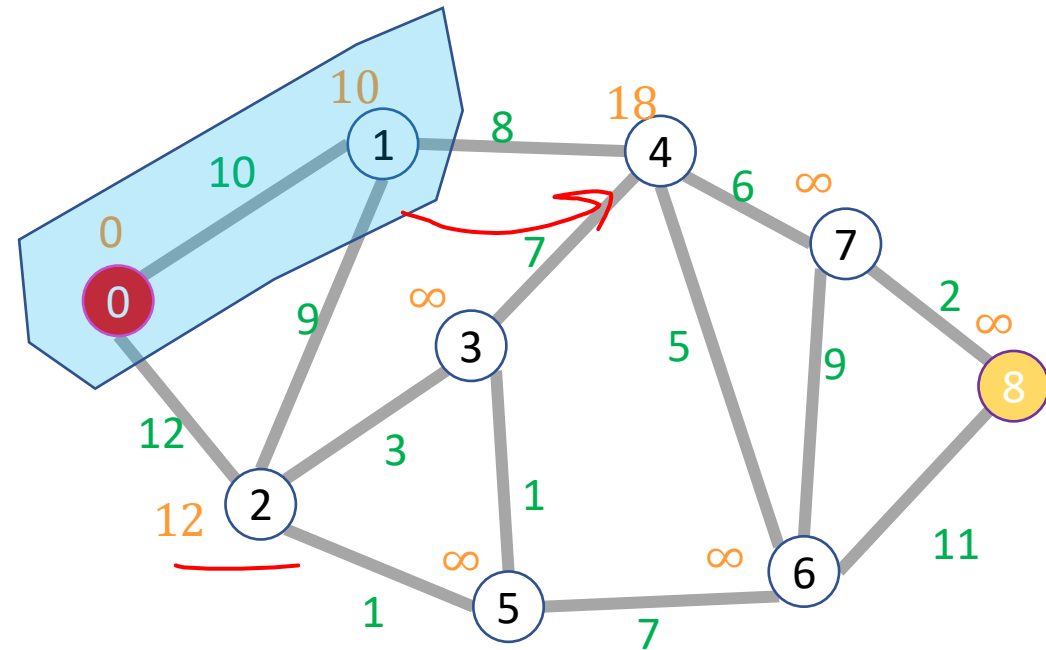| Node	Seen?	Done?	Distance
0	T	T	0
1	T	T	10
2	T	F	12
3	F	F	∞
4	T	F	18
5	F	F	∞
6	F	F	∞
7	F	F	∞
8	F	F	∞

Idea: When a node is the closest not-done thing to the start, we have found its shortest path

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

Dijkstra's Algorithm

Start: 0

End: 8

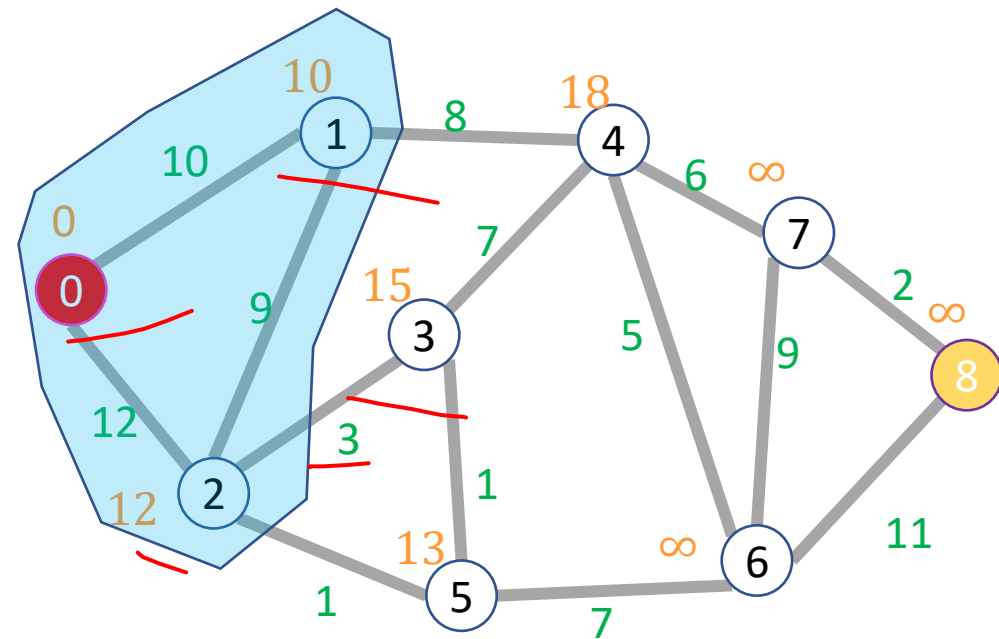| Node	Seen?	Done?	Distance
0	T	T	0
1	T	T	10
2	T	T	12
3	T	F	15
4	T	F	18
5	T	F	13
6	F	F	∞
7	F	F	∞
8	F	F	∞

Idea: When a node is the closest not-done thing to the start, we have found its shortest path

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

Dijkstra's Algorithm

Start: 0

End: 8

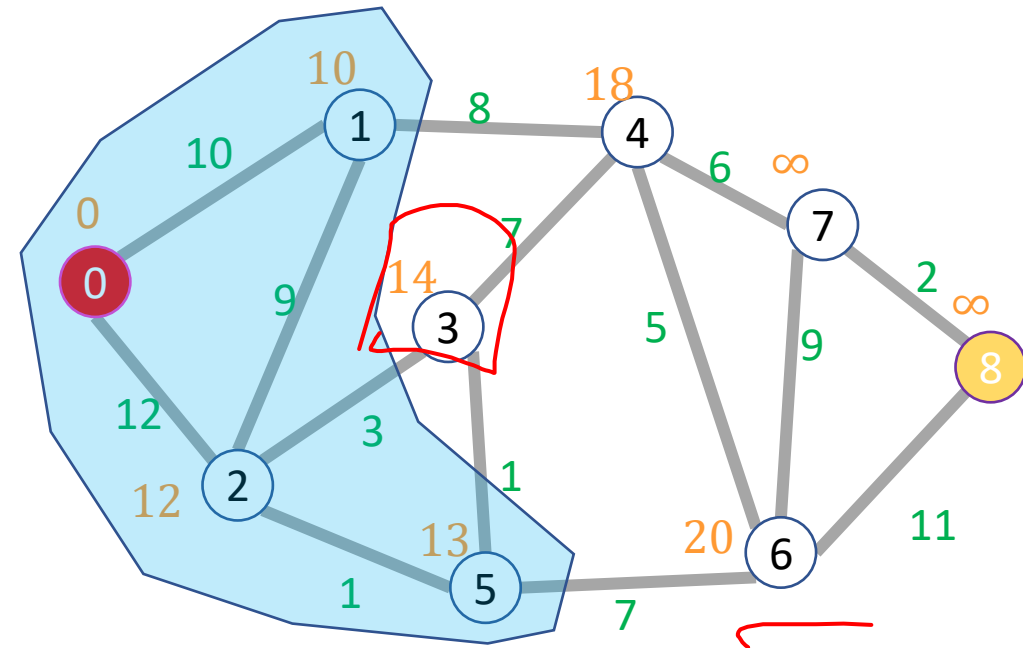| Node	Seen?	Done?	Distance
0	T	T	0
1	T	T	10
2	T	T	12
3	T	F	14
4	T	F	18
5	T	T	13
6	T	F	20
7	F	F	∞
8	F	F	∞

prev

Idea: When a node is the closest not-done thing to the start, we have found its shortest path

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

21

Dijkstra's Algorithm

Start: 0

End: 8

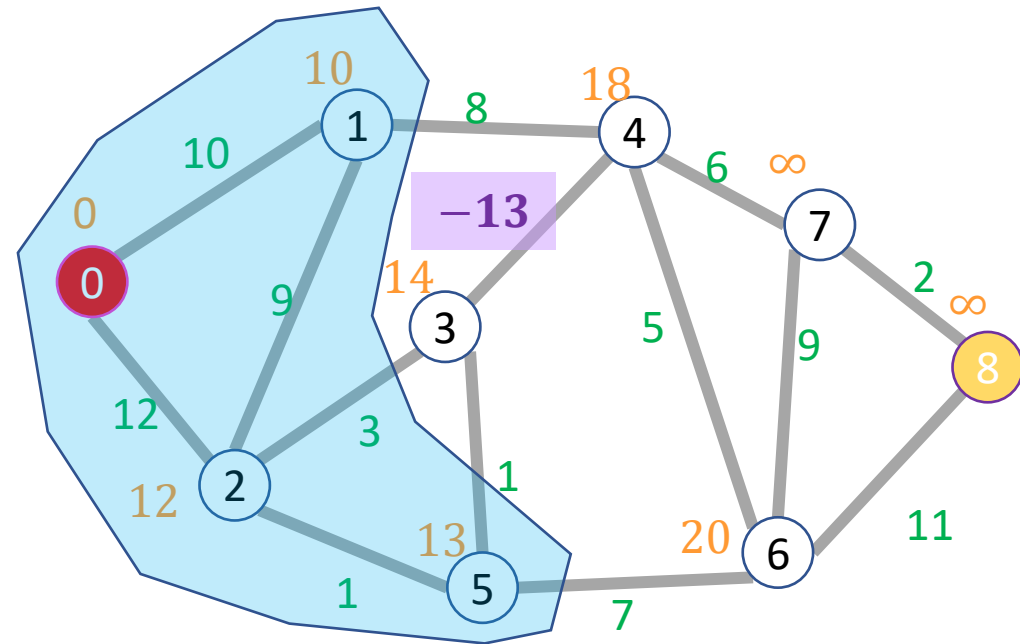| Node	Seen?	Done?	Distance
0	T	T	0
1	T	T	10
2	T	T	12
3	T	F	14
4	T	F	18
5	T	T	13
6	T	F	20
7	F	F	∞
8	F	F	∞

What if we had a negative-weight edge?

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

Dijkstra's Algorithm

Start: 0

End: 8

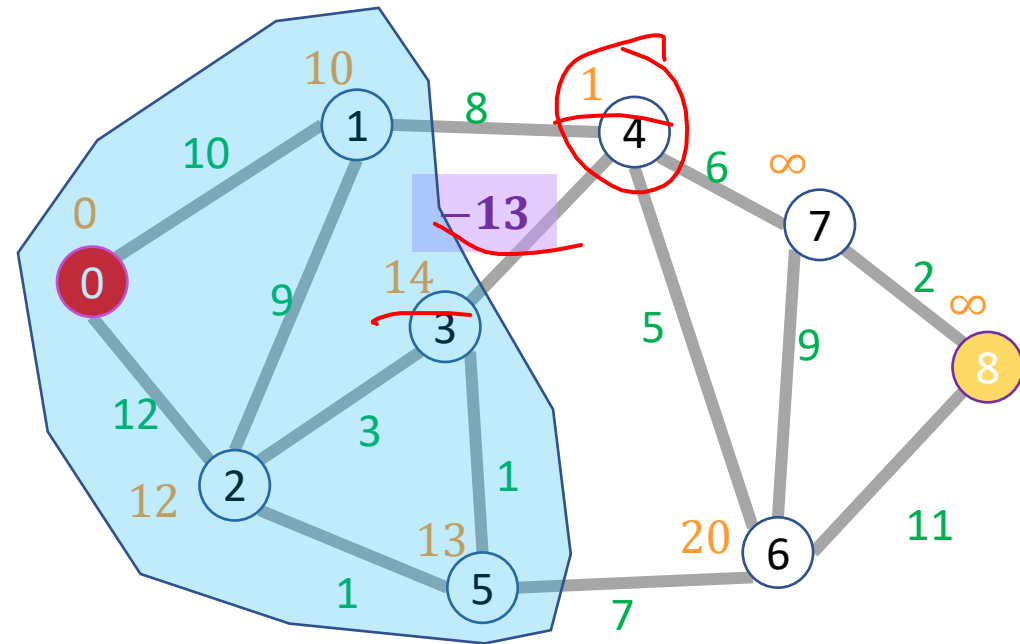| Node	Seen?	Done?	Distance
0	T	T	0
1	T	T	10
2	T	T	12
3	T	T	14
4	T	F	1
5	T	T	13
6	T	F	20
7	F	F	∞
8	F	F	∞

What if we had a negative-weight edge?

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
    Update its distance if we found a better path

Dijkstra's Algorithm

Start: 0

End: 8

Node	Seen?	Done?	Distance
0	T	T	0
1	T	T	10
2	T	T	12
3	T	T	14
4	T	T	1
5	T	T	13
6	T	F	20
7	F	F	∞
8	F	F	∞

There's a better path!

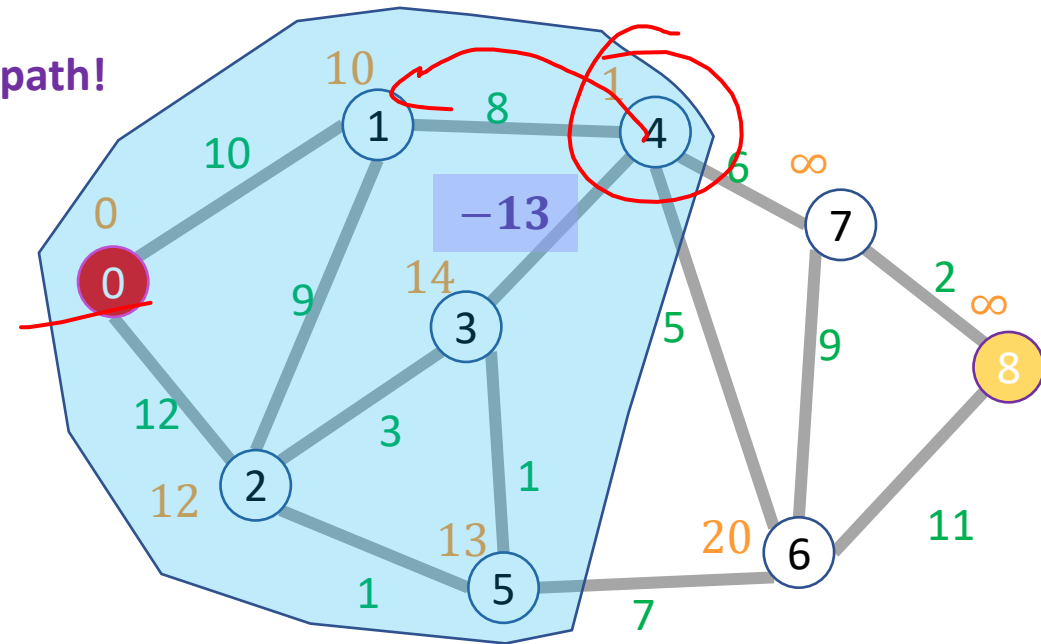What if we had a negative-weight edge?

Extract a node from priority queue (making it "done")
Mark extracted node as seen
for each not-done neighbor:
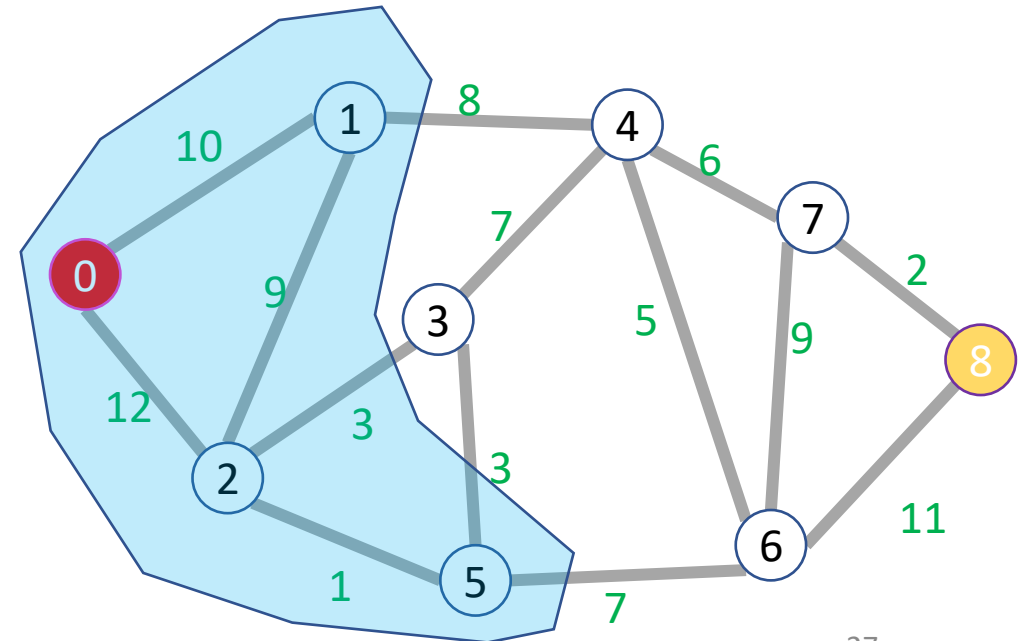   Update its distance if we found a better path

Dijkstra's Algorithm

```
int dijkstras(graph, start, end){
        distances = [∞, ∞, ∞,…];  // one index per node
        seen = [False,False,False,…];  // one index per node
        done = [False,False,False,…];  // one index per node
        PQ = new MinHeap();
        PQ.insert(0, start);  // priority=0, value=start
        distances[start] = 0;
        while (!PQ.isEmpty){
                current = PQ.extract();
                done[current] = True;
                for (neighbor : current.neighbors){
                        new_dist = distances[current]+weight(current,neighbor);
                        if (! seen[neighbor]){
                                seen[neighbor] = True;
                                distances[neighbor] = new_dist;
                                PQ.insert(new_dist, neighbor);
                        }
                        else if (! done[neighbor] && new_dist < distances[neighbor]){
                                distances[neighbor] = new_dist;
                                PQ.decreaseKey(new_dist,neighbor);
                        }
                }
        }
        return distances[end]
}
```

Dijkstra's Algorithm: Running Time

- How many total priority queue operations are necessary?
  - How many times is each node added to the priority queue?
  - How many times might a node's priority be changed?
- What's the running time of each priority queue operation?
- Overall running time:
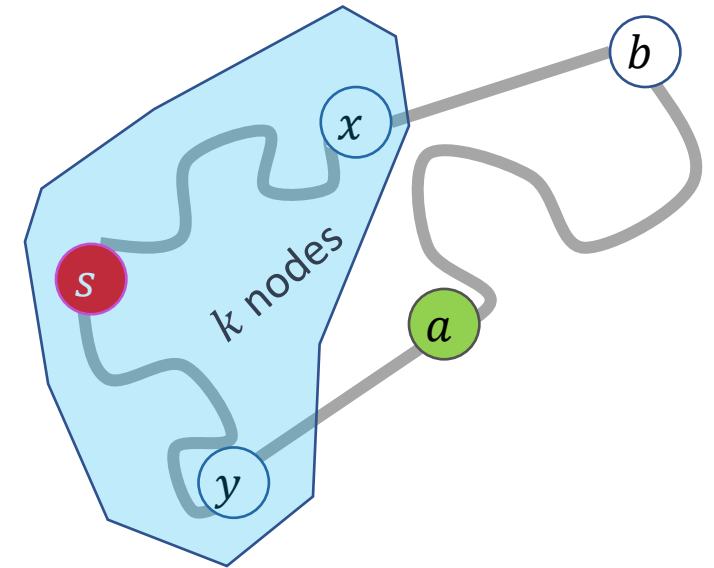 - $\Theta(|E|\log|V|)$

Dijkstra's Algorithm: Correctness

- Claim: when a node is removed from the priority queue, we have found its shortest path

- Induction over number of completed nodes

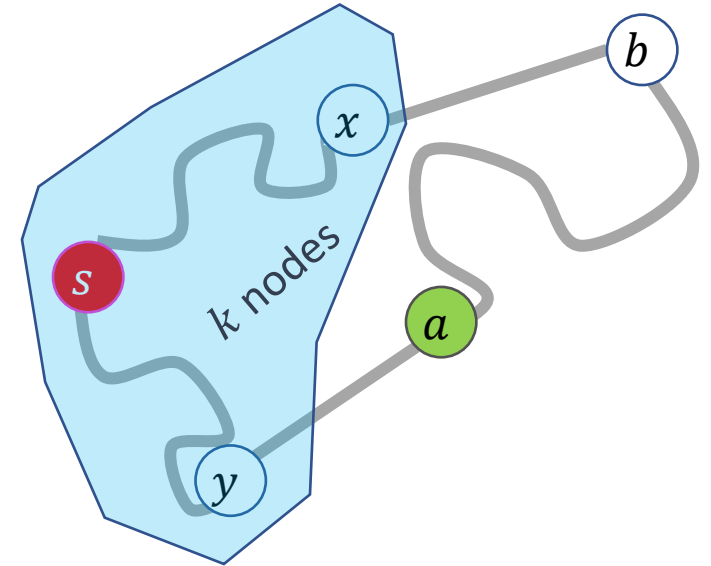- Base Case:

- Inductive Step:

Dijkstra's Algorithm: Correctness

- Claim: when a node is removed from the priority queue, its distance is that of the shortest path

- Induction over number of completed nodes

- Base Case: Only the start node removed
  - It is indeed 0 away from itself

- Inductive Step:
  - If we have correctly found shortest paths for the first k nodes, then when we remove node $k + 1$ we have found its shortest path
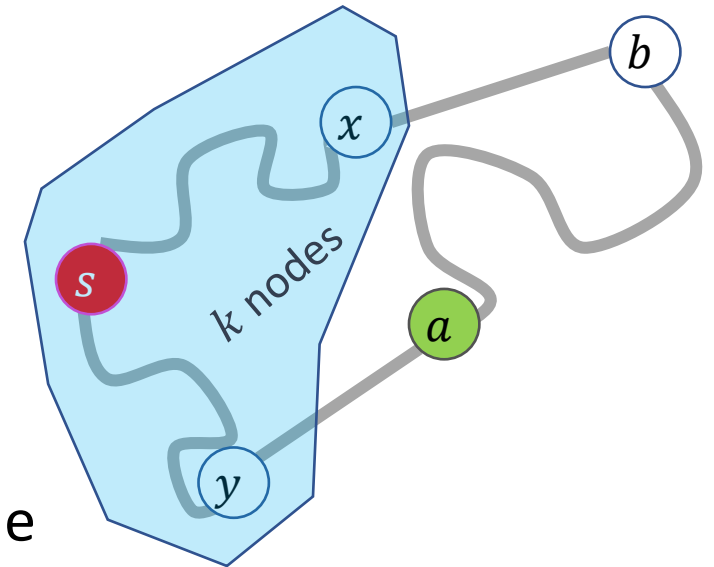
Dijkstra's Algorithm: Correctness

- Suppose a is the next node removed from the queue. What do we know bout a?

Dijkstra's Algorithm: Correctness

- Suppose a is the next node removed from the queue.
  - No other node incomplete node has a shorter path discovered so far

- Claim: no undiscovered path to a could be shorter
  - Consider any other incomplete node b that is 1 edge away from a complete node
  - a is the closest node that is one away from a complete node
  - Thus no path that includes b can be a shorter path to a
  - Therefore the shortest path to a must use only complete nodes, and therefore we have found it already!

Dijkstra's Algorithm: Correctness

- Suppose a is the next node removed from the queue.
  - No other node incomplete node has a shorter path discovered so far

- Claim: no undiscovered path to a could be shorter
  - Consider any other incomplete node b that is 1 edge away from a complete node
  - a is the closest node that is one away from a complete node
  - No path from b to a can have negative weight
  - Thus no path that includes b can be a shorter path to a
  - Therefore the shortest path to a must use only complete nodes, and therefore we have found it already!